

CS409 : Neural Networks (Semester II - 2021/22)

Unit 6: Convolutional Neural Networks (CNNs) (1)

Dr. Ruwan Nawarathna
Department of Statistics & Computer Science
Faculty of Science
University of Peradeniya

Content


- Introduction to Convolutional Neural Networks (CNNs)
- What is Convolution?
- Types of layers of a CNN

Convolutional Neural Networks (CNNs)

- **Convolutional networks**, also known as **convolutional neural networks** or CNNs, are a specialized kind of neural network for processing data that has a known, grid-like topology.
- Examples include:
 - image data, which can be thought of as a 2D grid of pixels.
 - time-series data, which can be thought of as a 1D grid taking samples at regular time intervals.


CNNs

- Convolutional networks have been tremendously successful in practical applications.
 - Computer vision
 - Object classification and detection in photographs
 - Natural language processing
 - Speech recognition



CNNs

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that does convolutional operation.


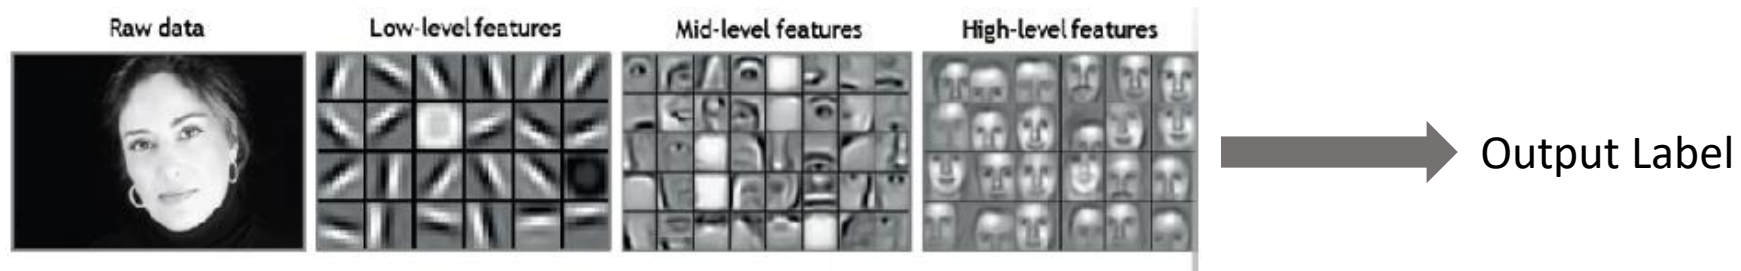


CNNs

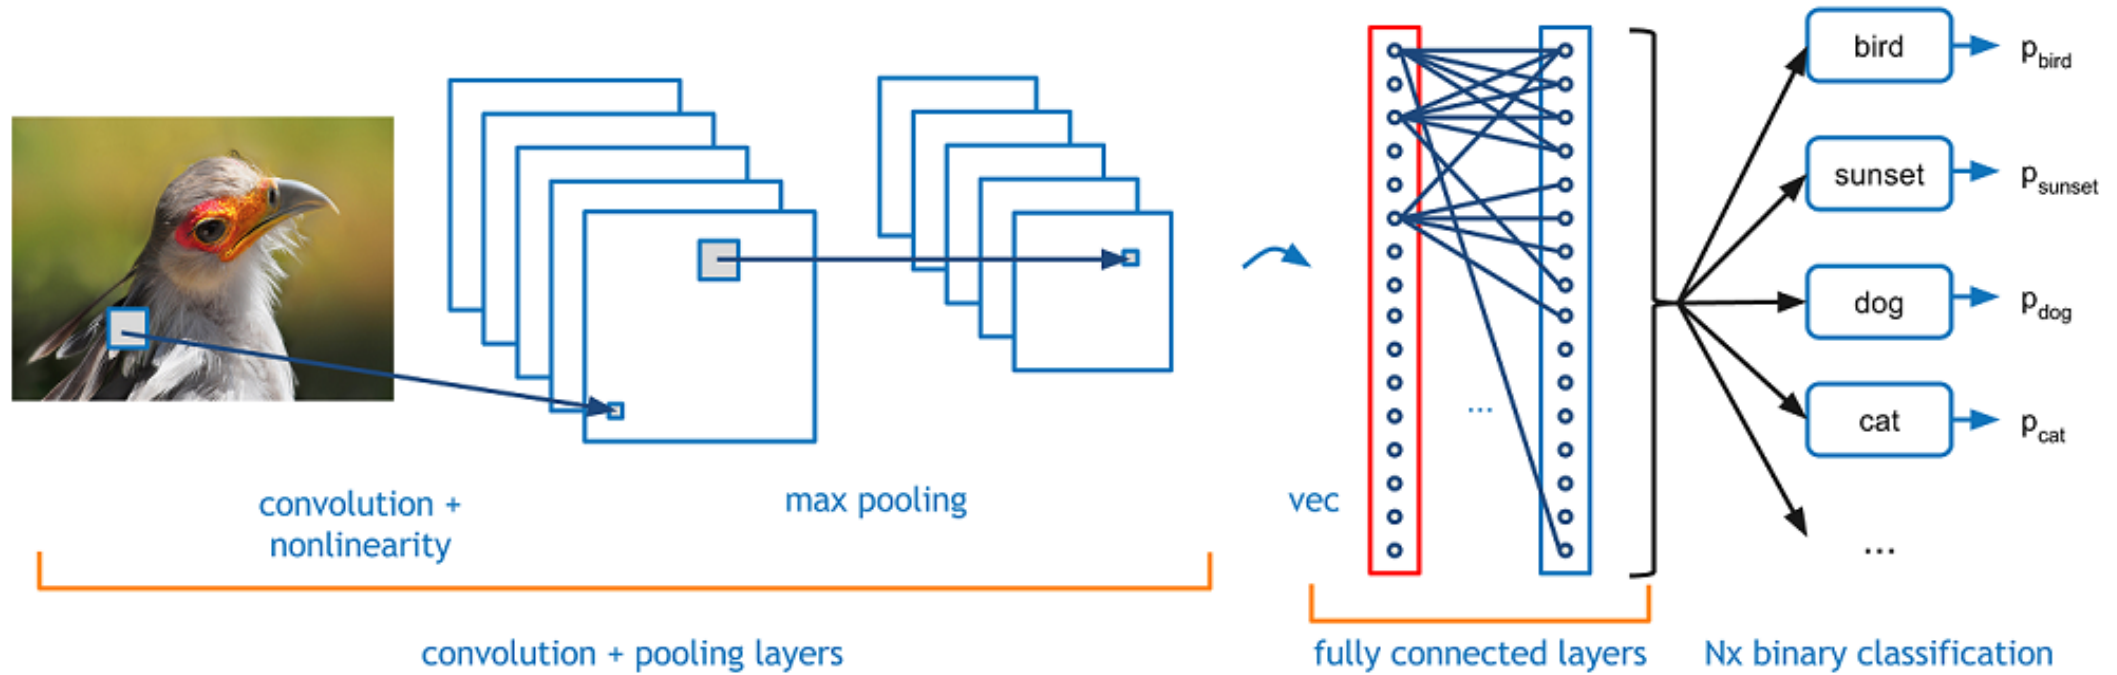
- The name “convolutional neural network” indicates that the network employs a mathematical operation called **convolution**.
- Convolution is a specialized kind of linear operation.
- Convolutional networks are simply neural networks that use convolution in place of general matrix multiplication in at least one of their layers.

CNNs

- Through series of convolutions (convolutional layers) feature learning is performed at various levels.



CNNs



Types of layers

- Three main types of layers are used to build CNN architectures:
 1. Convolutional layers (CONV)
 - Output: Feature Map
 - ReLU (Rectified linear unit) layers (RELU)
 2. Pooling (or Subsampling) layers (POOL)
 3. Fully connected layers (classification) (FC)
 - Multi-layer perceptron


Convolutional layer

- Rectangular grid of neurons
- Input from a rectangular section of previous layer
- Weights are same for each neuron
- Weights specify convolutional filters
- Several grids in each layer, each grid takes inputs from all layers using different filters

Convolution

- Convolution is a common image processing technique that changes the intensities of a pixel to reflect the intensities of the surrounding pixels.
- A common use of convolution is to create image filters

Convolution



Effect of Convolution



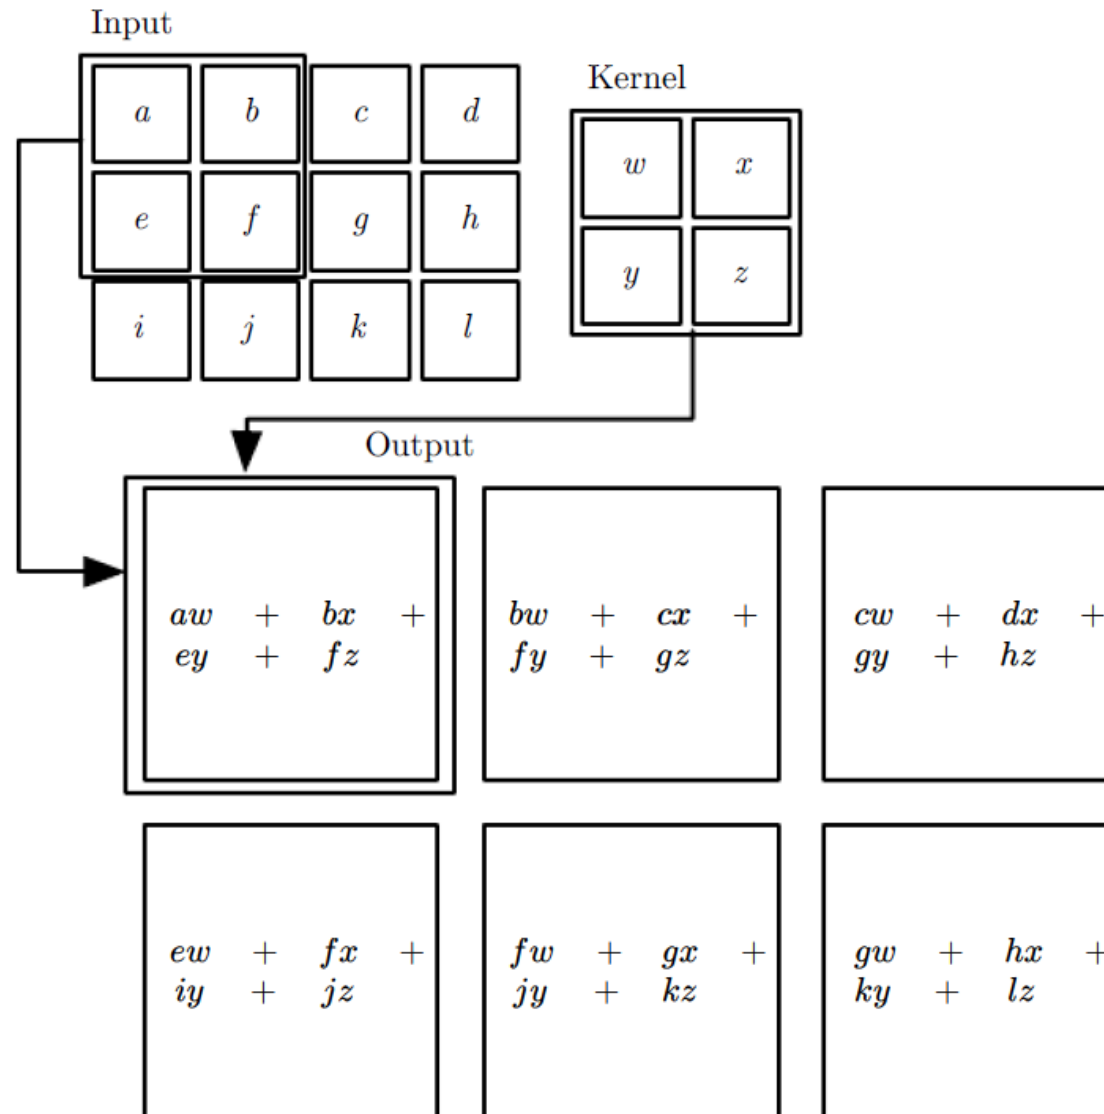
Original Image



Convolved image

This is an output of smoothing filter

Convolution



Convolution

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

These are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮ ⋮

Each filter detects a small pattern (3 x 3).

Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Dot
product



Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Repeat this for each filter

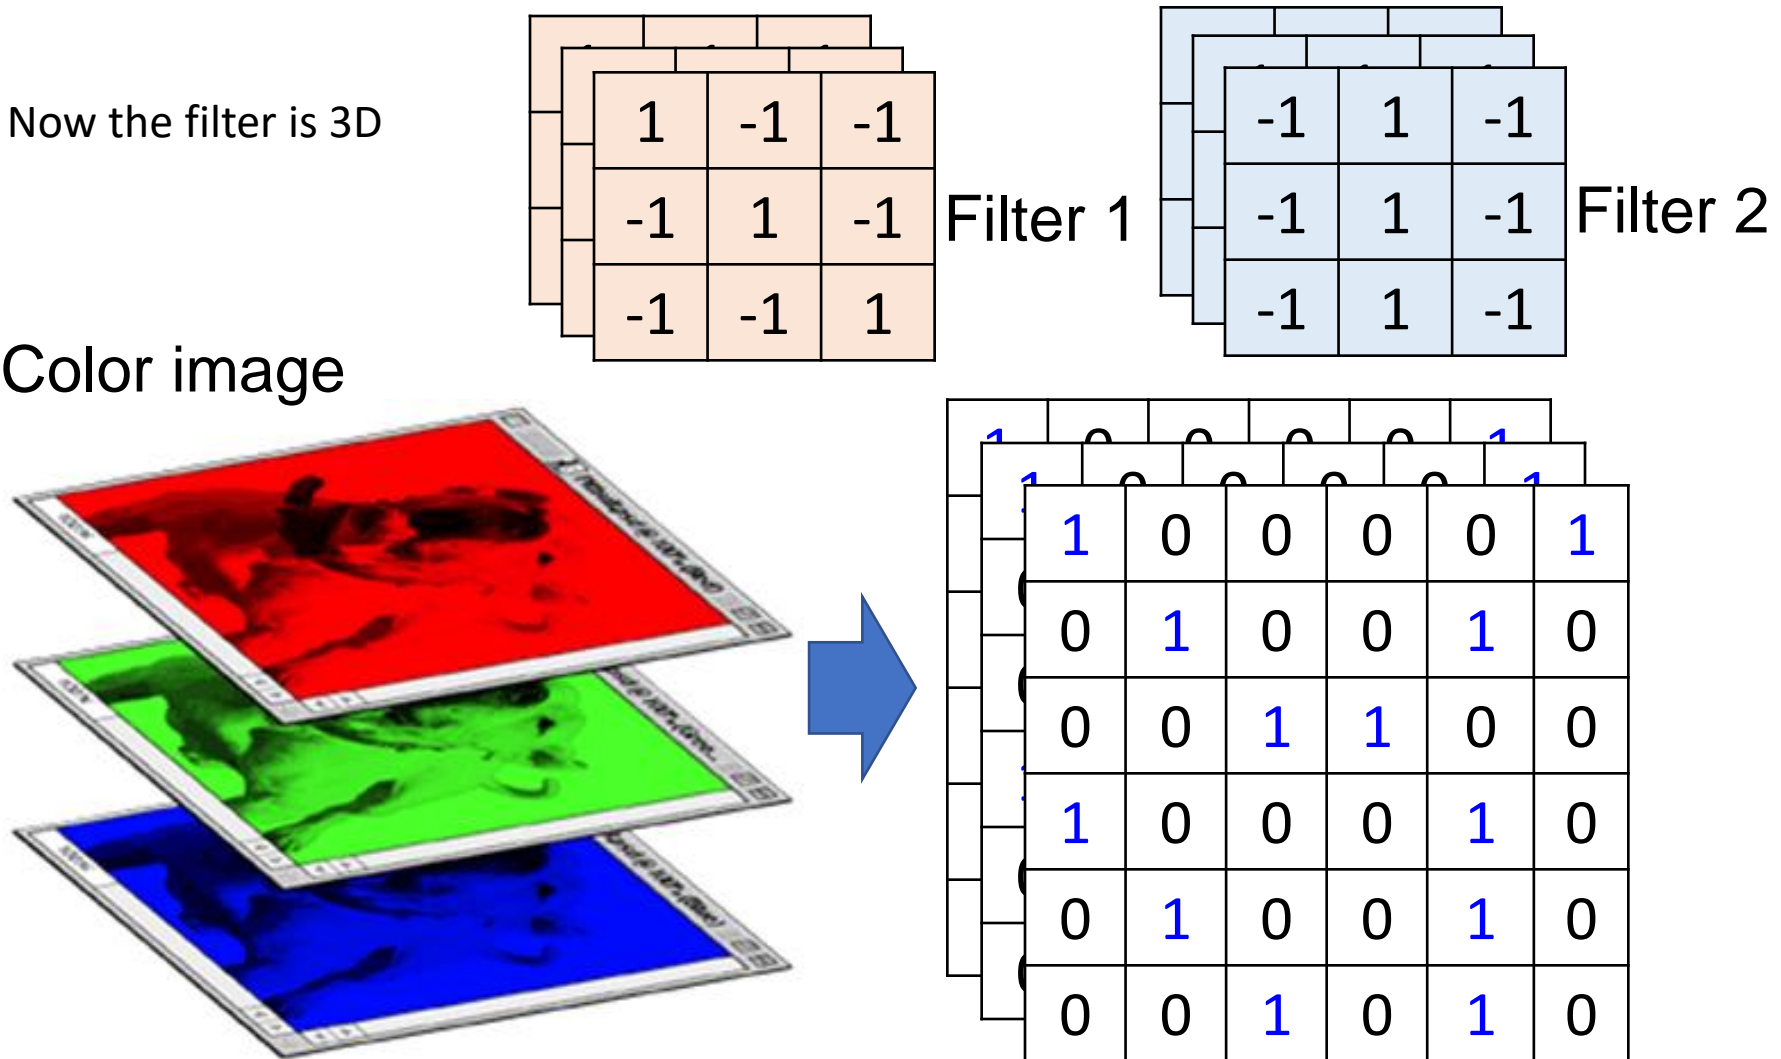


Two 4 x 4 images
Forming 2 x 4 x 4 matrix

Convolution on Color images: RGB 3 channels

Now the filter is 3D

Color image



Convolution with Neurons

weights

1	-1	-1
-1	1	-1
-1	-1	1

Convolution

$$z = 1 \times 1 + -1 \times 0 + -1 \times 0 + -1 \times 0 + 1 \times 1 + -1 \times 0 + -1 \times 0 + -1 \times 0 + 1 \times 1$$
$$= 3$$

Inputs

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

1

-1

-1

-1

1

-1

-1

-1

1

Neuron

z

ReLU

y


$$z = 1 \times 1 + -1 \times 0 + -1 \times 0 + -1 \times 0 + 1 \times 1$$
$$+ -1 \times 0 + -1 \times 0 + -1 \times 0 + 1 \times 1$$
$$= 3$$

$$y = \text{ReLU}(3) = 3$$


There could
be a bias
weight

We get same output as with convolution


Convolution with Neurons



Convolution with Neurons – 3D Inputs




Convolutional Layers : Parameter Sharing and Local Connectivity




- Parameter sharing is sharing of weights by all neurons in a particular feature map.
- Local connectivity is the concept of each neuron connected only to a subset of the input image (unlike a neural network where all the neurons are fully connected)


Parameter Sharing and Local Connectivity



fewer parameters!




Parameter Sharing and Local Connectivity




Fewer parameters

Even fewer parameters




Zero Padding

- Adding zero padding will preserve the size spatially.
- In general, common to see CONV layers with stride 1, filter size $F \times F$, and zero-padding with $(F-1)/2$.




Stride

$S = 1$




$S = 2$



Convolution Layer – Size of Activation Maps (Output Volume Size)

32x32x3 image




5x5x3 filter



Since the image is 3-dimensional, the filter is also 3-dimensional


That is depth of the input = depth of the filter

Convolution Layer – Size of Activation Maps (Output Volume Size)



Convolution Layer – Size of Activation Maps (Output Volume Size)


For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:




We stack these up to get a “new image” of size 28x28x6!

Convolution Layer – Size of Activation Maps (Output Volume Size)

- CNN is a sequence of Convolutional Layers, interspersed with activation functions



Calculating the Output Volume Size and Other Parameters



Filter Extent = F
Stride = S
Padding = P

Produces a volume of size $W_2 \times H_2 \times D_2$ where:

$$W_2 = (W_1 - F + 2P) / S + 1$$


$$H_2 = (H_1 - F + 2P) / S + 1$$

$$D_2 = K$$

With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and k biases

Convolution Layer – Size of Activation Maps (Output Volume Size)

- Example



Input Volume of 32x32x3

Stride = 1

Padding = 2

Apply 10 5x5 filters

$$(W_2=H_2) = \frac{32-5+2*2}{1}+1$$
$$= 32$$


$$D_2 = 10$$

Output Volume Size = 32 x 32 x 10


Number of shared parameters (weights) per filter
= $5 \times 5 \times 3 + 1(\text{bias}) = 76$

Total number of shared parameters (weights)
= $5 \times 5 \times 3 + 1(\text{bias}) = 76 \times 10 = 760$

Example CNN – with CONV Layers Only



Example CNN – with CONV Layers Only



Same network with different Representation (AlexNet Style)

Pooling Layers

- Why Pooling: Subsampling pixels will not change the object

bird



Subsampling
(Pooling)


bird



We can subsample the pixels to make image smaller
fewer parameters to characterize the image

Pooling Layer

- makes the representations smaller and more manageable



Pooling layer

- Takes smaller blocks from convolutional layer
- Subsamples to produce single output from that block
- Several ways- average or maximum or learned linear combination of neurons
- For example, max pooling layers take maximum out of that block
- Pooling layers do not have trainable weights.

Types of Pooling

- Max pooling
- Average pooling
- Global max pooling
- Global average pooling

Pooling Layer

- Max Pooling

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

max pool with 2x2 filters
and stride 2



6	8
3	4

- Average Pooling

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

Avg pool with 2x2 filters
and stride 2



3.25	5.25
2	2

Effect of Pooling




Original Image



Convolved image




Convolved image

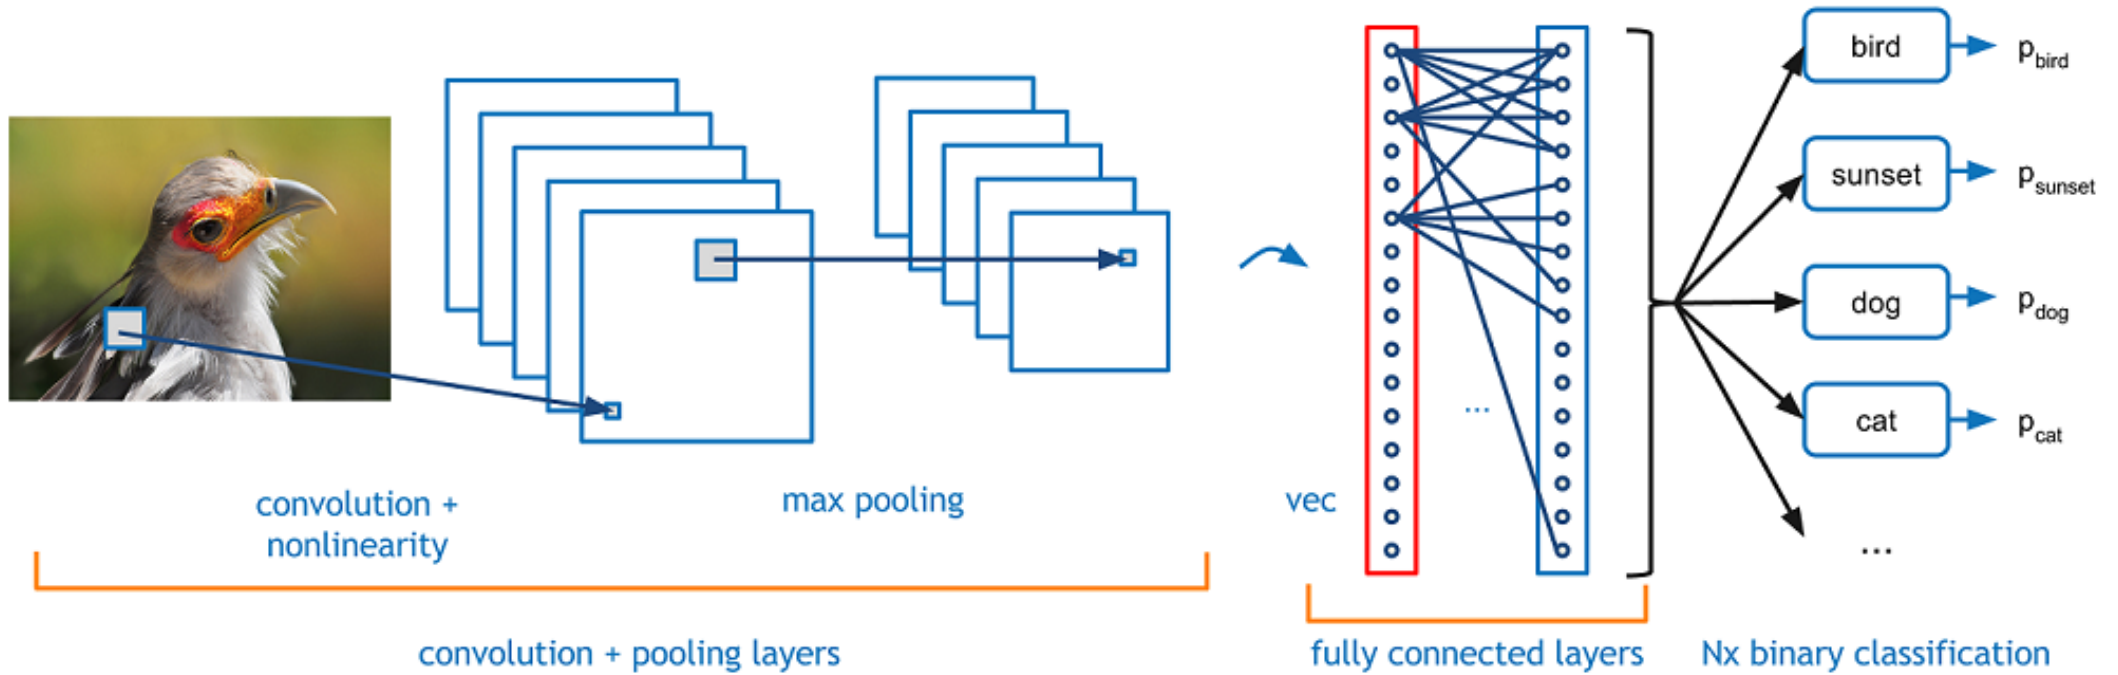


MAX Pooling

Example CNN – with CONV and Pooling Layers




Fully-connected layer




Fully-connected layer

- Finally, after several convolutional and max pooling layers, the high-level reasoning in the neural network is done via fully connected layers.
- Neurons in a fully connected layer have full connections to all activations in the previous layer, as seen in regular Neural Networks.

Example CNN




Fully-connected layer



Last output volume will be flattened and use it for the input layer of the fully connected network.

In this case, flattening gives $64 \times 4 \times 4 = 1024$ Input values which are the outputs of 1024 neurons of the output volume



Example Architecture – LeNet 5

- Exercise: Explain the architecture of LeNet-5

